

SQL



Relational Data Model

- Pros**
- > Easy to use and setup.
 - > Universal, compatible with many tools.
 - > Good at high-performance workloads.
 - > Good at structure data.
- Cons**
- > Time consuming to understand and design the structure of the database.
 - > Can be difficult to scale.

No SQL



Document Data Model

- Pros**
- > No investment to design model.
 - > Rapid development cycles.
 - > In general faster than SQL.
 - > Runs well on the cloud.
- Cons**
- > Unsuitable for interconnected data.
 - > Technology still maturing.
 - > Can have slower response time.

NOSQL SUMMARY

MongoDB and HBase

Relational (SQL)	MongoDB
Database	Database
Table	Collection
Row	Json documents
Index	Index
Join	Embedded document

- database can have one or more collections
- collection, you can have none (empty collection), one or more JSON documents

show dbs	show the databases
use my_database	create or connect to a database
db	show the database I am using
show collections	shows the collections in the current database
db.my_collection.insert ({ attribute1:"hello", attribute2: 100 })	insert a document in a collection
db.other_collection.find ()	
db.other_collection.find().pretty()	
db.other_collection.find().count()	
db.other_collection.find().pretty().count()	
db.agenda.find({age: 25})	filter the documents with a condition
db.agenda.find ({ "address.province": "Murcia" })	filter the documents with a condition with Embedded Document
db.agenda.find ({children: "Ines" })	filter the documents with a condition with array field

\$gt or \$gte	Greater than or equal age: {\$gt:25}
\$lt or \$lte	Lower than or equal
\$eq or \$ne	Equal or not equal
\$in or \$ni	In or not in

{name: /^Juan/ }	Begins with Juan
{name: /Juan\$/ }	End with Juan
{name: /Juan/ }	Contains Juan
{name: /Juani/ }	Contains Juan(case insensitive)

db.agenda.find ({ \$or: [{ "address.province":"Madrid" }, { "age":{\$gte:30} }] })	or
--	----

db.agenda.find ({ \$and:[{age:{\$gte:30}} , {"address.province":"Madrid"}] }, { name:1, lastName:1, _id:0 })	And (when using the same field many times in the conditions.)
db.agenda.find({... }, {name : 1 , " _id" : 0 })	Filtering field
db.test.remove({}) or db.test4.deleteMany({ y:2 })	Remove all documents of test collection (don't forget : {})
db.test2.remove({ y:2 } , {justOne: true}) or db.test2.deleteOne({ y:2 })	Delete just 1 doc
db.test.drop()	Dropping a collection
db.<collection>.update ({ <conditions> }, {<modification> }, {<options> }) db.test5.update ({ _id:1}, {y:12}) db.test6.update ({x:1}, {\$set:{x:100}}, {multi:true, upsert:true})	Updating document You can also add a field that did not exist Multi : multiple updates (only partial update (\$) Upsert : conditions does not match any document a new document is inserted instead works with partial and whole

\$set:{}	Sets the value of a field in a document
\$push:{}	Add value to an Array
\$inc:{}	Increments the value of the field by the specified amount
\$mul:{}	Multiplies the value of the field by the specified amount
\$rename:{x : "z"}	Renames a field in a document
\$unset:{}	Removes the specified field from a document.
\$pull:{}	Remove an element from an array

db.test6.update ({x:1},{\$set:{x:100}})
id field cannot be updated > get an error

db.test.insert ({name: "Luis", age: null })	In lowercase A string with the word null, isn't a null value. No value will also show up if age: null
db.test.find({ age: { \$exists: true } }, { _id: 0 })	Show the documents that have the age field
db.test.find ({ age: { \$exists: false } }, { _id: 0 })	Show the documents that do not have the age field
db.airplanes.find().sort({"passenger.name":1})	To order by (1 : ascending, -1 : descending)
db.tickets.find().limit(10)	Show only 10 doc

db.<collection>.help()	
db.countries.distinct ("field", {"condition"})	Find the the different values for a field.
db.countries.distinct ("name", { continent: "AMERICA" })	Find the different countries for the continent America.
db.countries.distinct ("name", { \$and: [{continent: "AMERICA"},{population: {\$gt:35}}] })	Find the different countries for the continent America with population greater than 35 millions
db.countries.aggregate([{ \$group : { . . . } }])	Aggregate Method for Grouping in MongoDB
db.countries.aggregate ([{ \$group : { _id: "\$continent", number_of_countries: {\$sum: 1}, total_population: {\$sum: "\$population"}} }])	
db.countries.aggregate([{\$group : { _id: "\$continent", max_population: {\$max: "\$population"}, avg_population: {\$avg: "\$population"}, min_population: {\$min: "\$population"} } }])	
db.countries.aggregate([{\$group : { _id : null, total_population:{\$sum:"\$population"} } }])	_id is null everything is just 1 group SELECT SUM(POPULATION)FROM COUNTRIES
db.countries.aggregate([{ \$match : {<conditions>} }, { \$group : { . . . } }])	Condition can be established before doing the groups

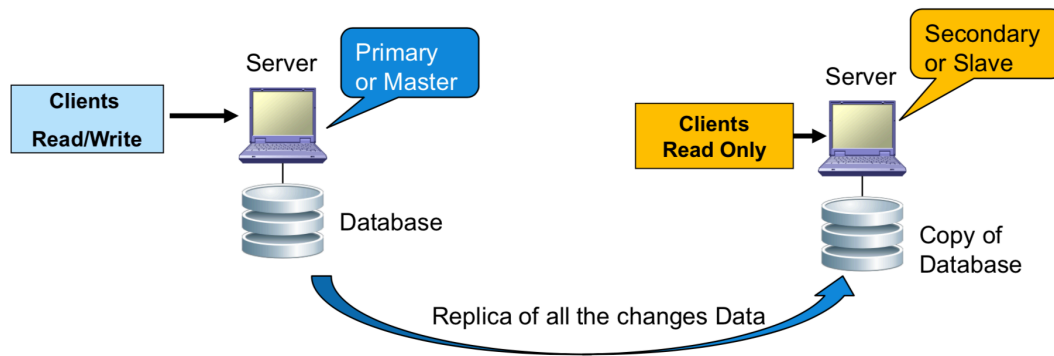
\$sum	Returns a sum of numerical values. Ignores non-numeric values.
\$avg	Returns an average of numerical values. Ignores non-numeric values.
\$max	Returns the highest expression value for each group.
\$min	Returns the lowest expression value for each group.
\$first	Returns a value from the first document for each group.

\$last	Returns a value from the last document for each group.
\$push	Returns an array of expression values for each group. (an array with all the values of the field.)


db.test.find ({phrase:/tailor/})	Regular Expressions
OR	
db.test.find ({phrase:{\$regex:"tailor"}})	
i	Case Insensitive
^	Beginning of a string
\$	End of a string
\s	whitespace
\S	anything but whitespace
\d	1 digit (Find the documents that are only numbers)
\D	1 character that is not a digit (Find the documents that at least have 1 letter (or blank)
m[aeiou]	Find the String that contains ma, me, mi, mo, mu
mother kitchen	Find the word mother or kitchen
\	escape character
db.test.find ({x:/\\$/})	Search for the documents containing \$

Replication

- Make a **copy** (replica) of all the data in another server.
- When a **change** occur in the database the same change is propagated to the replica.
- The replication is **asynchronous** (it takes some time)



- When the Primary **fails** the secondary server becomes primary.
- Clients reconnect to secondary
- When failed Primary is back to life it can become Secondary.
- Receives the changes from the new primary server.



- Multiples slaves : MongoDB
- Every server primary: IBM Cloudant

Sharding

To increase computer power I can do it 2 ways:

- **Vertically**: Adding more CPU, Memory or disk to the same server
- **Horizontally**: Adding more Servers and dividing the documents among

Column Family: HBase (columns family)

Facebook uses it for messaging (Powered by HBASE)

- History
 - o Google created in 2004 a new database called Big Table.
 - o Google published "Big Table" paper in 2006:

- “Bigtable: A Distributed Storage System for Structured Data”.
- HBase (and Casandra) are based in **Big Table**:
 - Manage high volume of data.
- Hbase(and Casandra) are **open source software**.
 - Hbase (and Casandra) are **Apache Top Level Projects**:
<http://hbase.apache.org/>
- HBase Table
 - The basic element in Hbase
 - With Columns and Rows (Similar to Relational).
 - The Hbase Table is a bit different from a relational database table:
 - Even It has rows and columns as well.
 - It has other components like new **Column-family** concept.
 - You can **not** use SQL to access it (NoSQL).
 - You can **not** do Join (only 1 table at the time).
 - Similar than with MongoDB Collections.
 - Resolve the Join your self programming.

Important HBase Concept

- **Column Family**
 - A group of columns.
 - A table can have several columns familiy.
 - A column is represented by the
 - Column family: Column name:
 - Personal_data:Name
 - Personal_data:Age
 - Internet_data:Twitter
 - In RDBMS there are no column-families.

Table Persons:

Column-family	➡	Personal_data			Internet_data	
Columns	➡	Name	Address	Age	Twitter	E-mail
Rows	{	John	White	25	@lg	lg@x.com
		Will	Smith	30	@wsmith	w@y.es
		Bruce	Lee	35	@lee	lee@p.tv

- **Row Key**
 - Every row in the table has a unique value that identifies one row from the others of rows.
 - Rows are stored sorted by Row Key.
 - When inserting a row the user/applications must give a value for the row key.
 - Similar to the primary key/unique key concept in RDBMS.
 - Flexible schema:

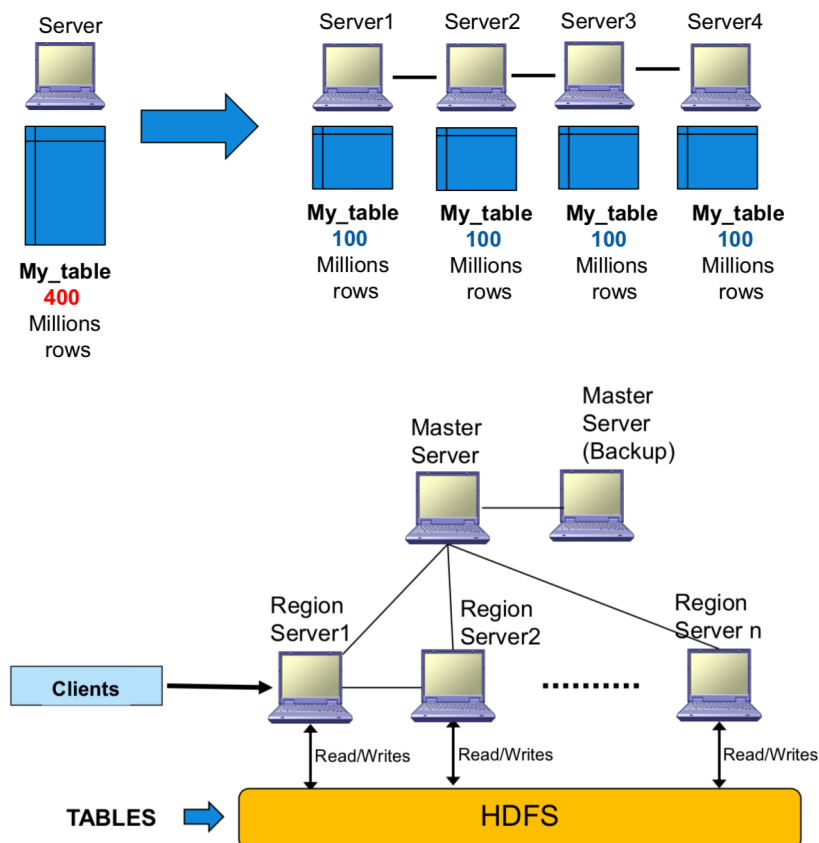
- Columns are “created” when inserting data in a row.
 - When creating a table you only give the name of the table and the column families names, not the columns.
 - Column names can be the same in different column families.
 - **Null values does not exist** => nothing is stored for a missing column (“sparse data”).
- **Cell and Timestamp**
- Timestamp=Time when the value was inserted or modified.
 - Everytime the value changes the timestamp is inserted.
 - It is inserted by Hbase but you can insert instead.
 - **Versioning**
 - Keeping past values of the data: When application do an update I keep the old value before the update.
 - By default Hbase keeps only 1 version but it can be customized as you wish.
 - In older versions of Hbase (until v0,96) default was 3 values.
 - When quering data if not specified I will obtain the last value but I can ask for a past value.
 - The timestamp value mentioned below is the one used for versioning:
 - The highest timestamp is the last value.
 - Using this timestamp I can obtain a past value.
 - The timestamp is a long integer in UTC format: Number of milliseconds since midnight, January1, 1970 UTC

Row Key	Personal_data			
1000	Name	Timestamp	Age	Timestamp
	Luis Reina	1433163936648	35	1522163936635
1000	Name	Timestamp	Age	Timestamp
	Luis Reina	1433163936648	25	1522163936875
1000	Name	Timestamp	Age	Timestamp
	Luis Reina	1433163936648	30	1522163936999

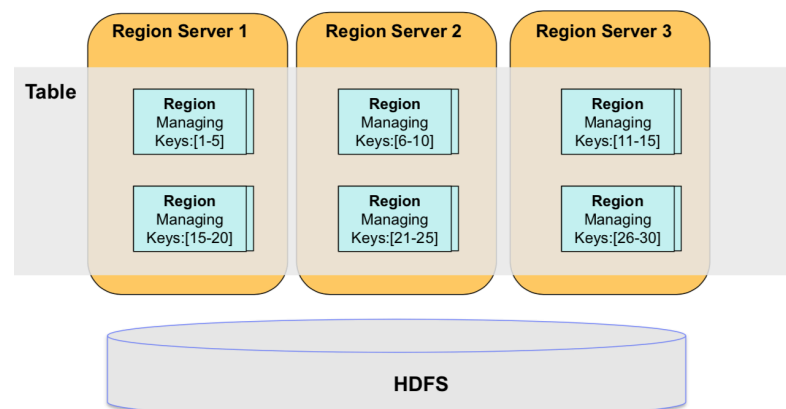
HBase Architecture

- HBase uses HDFS from Hadoop: Data in HBASE is **stored** in **HDFS**.
 - The standalone allows to run out of HDFS but is only for development never for a production system.
- This way data is **replicated** among different servers
 - In MongoDB I need to specify the replicas.
 - In Hbase is done by HDFS **automatically**: Thanks to HDFS redundancy of the blocks.
- Because of this Hbase is recognized as the **Hadoop database**.

- It does not use **Map/Reduce**.
- Map/Reduce: Batch reads (slow for queries accessing little data).
 - o Quick random Access (Real Time) to the data (read and writes).
- **HBase Sharding**
 - o As mentioned sharding allows to split the data among different servers to scale horizontally (adding servers):
 - In MongoDB we Split the **json documents** among the different servers
 - In Hbase we splits the **rows** among the servers.




- **Region Servers and Regions**
 - o HBase Cluster is a number of Region Servers.
 - o Every Region Server has 1 or more Regions.
 - o Data is divided (sharded) in the different Regions.



HBASE SHELL

- The Hbase Shell is a client where we can write the HBase commands to interact with an Hbase Server.
 - o Remember the Client/server Architecture:



- use the AWS Cluster where Hbase is installed.

CREATING HBASE TABLES

- TO OPEN THE SHELL:
 - o Open a ssh connection to the Linux Server (Edge node 1) Connect to the Linux Edge Node using ssh client.
 - o Execute in Linux this command to open the HBase Shell
 - hbase shell

create 'Table_name', 'Column_family1', 'Column_family2'	creating a hbase table
or	
create 'Table_name', [{NAME => 'Column_family1'}, {NAME => 'Column_family2'}]	
list	List all tables
exists 'Table_name'	see if a table exists
Alter 'Table_name', 'Column_family1'	Add a New Column Family
Disable 'Table_name'	Dropping a Column
Drop 'Table_Name'	
Describe 'Table_name'	Show the "structure" and Information of a Table
Alter 'Table_name', {NAME => 'Column_family1', VERSIONS => 3}	Changing the number of versions for a Column Family

HBASE NAMESPACE

- logical way of grouping tables:
 - o Similar to Relational Schema.
 - o A way to group tables together
 - o For Namespace=jonas
 - jonas:table1
 - jonas:table2
 - o For Namespace=John

- John:table1
- ➔ Same Table name in a different Namespace

<code>creating_namespace 'Jonas'</code>	create the Namespace first
<code>create 'namespace : table_name' , 'column_families'</code> <code>create 'Luis:Table1', 'col_fam1'</code>	Then create the table under the Namespace
List_namespace	Listing namespaces There are 2 predefined namespaces: ✓ hbase: system namespace, used to contain HBase internal tables. ✓ default: tables created with no explicit specified namespace.
<code>Drop_namespace 'namespace'</code> <code>Drop_namespace 'Jonas'</code>	Dropping a namespace ✓ You can drop a namespace if it doesn't contain any table in it. You must first drop all tables and then the namespace.

INSERT DATA INTO HBASE TABLES

<code>put 'table_name', 'row_key', 'column_family:column', 'value'</code> <code>put 'my_table', 1, 'cf1:c1', 'hello'</code>	insert data
<code>put 'table_reina', 1, 'private:name', 'Luis Reina'</code> <code>put 'table_reina', 1, 'private:age', 30</code> <code>put 'table_reina', 1, 'public:email', 'lreina@faculty.ie.edu'</code>	Row key 1 column name value your name and age column your age both columns in column family private and column email with value your_email in column family public
<code>put 'table_reina', 2, 'private:name', 'Griezmann'</code> <code>put 'table_reina', 2, 'private:nationality', 'french'</code> <code>put 'table_reina', 2, 'public:email', 'jsmith@faculty.ie.edu'</code>	Row key 2 column name value your_friend and nationality column his nationality in column family private and column email with value his_email in column family public

QUERYING DATA IN HBASE

scan 'my_table'	showing all the rows of a table
scan 'my_table ', {COLUMN => 'cf1' }	Specifying the column family required, option COLUMN:
scan 'my_table ', {COLUMN => 'cf1:c1' }	Specifying the column required, option COLUMN
scan 'my_table ', {COLUMNS => ['cf1','cf2'] }	Specifying several column families required, option COLUMNS
scan 'my_table ', {COLUMNS => ['cf1:c1','cf2:c1'] }	Specifying several columns required, option COLUMNS
scan 'my_table ', {COLUMNS => ['cf1:c1','cf2'] }	Specifying a mix of columns and column families, option COLUMNS:
scan 'my_table ', {LIMIT => 10}	Limiting the Number of Rows Returned, option LIMIT:
scan 'my_table ', {STARTROW => '2'}	Identifying the first row to return, option STARTROW
scan 'my_table ', { COLUMNS => ['cf1:c1','cf2:c1'], LIMIT => 1 , STARTROW => '2' }	Multiple Options together (separate by commas)
get '<table_name>' , '<row_key>'	get command to read 1 row only
get '<table_name>' , '<row_key>', '<column families or columns>'	Specifying the column families or columns I need
get 'my_table', 1, {COLUMNS => 'cf1:c10'}	Some options are also valid for get
scan 'table_reina', {LIMIT=>1, STARTROW=>'2'} OR get 'table_reina', '2'	Read the second row

UPDATING DATA IN HBASE

No specific command for updating a value.

As commented Hbase does versioning:

- Using the Timestamp when the data was inserted.
- Hbase will keep the old values of the data: if we specify versions bigger than 1 when creating the table.

Put 'my_table', 1, 'cf1:c1', 'hello' put 'my_table', 1, 'cf1:c1', 'new value'	same put command is used with same row key
---	--

RETRIEVING PAST DATA

create 'my_table', {NAME=> 'cf1',VERSIONS=>3}	The column family must have VERSIONING enabled if it is disabled by default
alter 'my_table',NAME => 'cf1', VERSIONS => 3	The column family must have VERSIONING enabled if it is disabled by default
get 'my_table', 1,{TIMESTAMP=> 1433701796440}	There is an option TIMESTAMP
put 'my_table', 1, 'cf1:c1', 'hello'	timestamp1
put 'my_table', 1, 'cf1:c1', 'new value'	timestamp2
get 'my_table',1,{TIMESTAMP=> timestamp1}	ask for past data specfing a past timestamp
scan 'my_table',{VERSIONS=>3}	Show all the versions
scan 'my_table', {COLUMNS=>'cf1:c1', TIMERANGE=>[1390651029025,1490651031644]}	Show the newest version for a Time Range
scan 'my_table', {COLUMNS=>'cf1:c1', TIMERANGE=>[1390651029025,1490651031644], VERSIONS=>2}	Show 2 versions for a Time Range